Dataset Acquisition

The CheXpert Plus dataset was obtained from Stanford University's AI in Medicine and Imaging (AIMI) portal, which provides access to the dataset under a data use agreement. The full image dataset, comprising approximately 720GB across five zipped PNG chunks, was transferred directly from Stanford AIMI's Azure Blob Storage to a personal Azure Blob Storage account using AzCopy, bypassing local storage entirely to stay within compute constraints. Each chunk was downloaded sequentially into a Google Colab runtime, filtered to retain only frontal-view images from female patients, and uploaded to the personal container before the chunk was deleted to manage disk space. The accompanying metadata CSV and CheXBERT label files were downloaded separately and merged to produce the working dataset. No modifications were made to the raw images. Label preprocessing involved extracting binary pleural effusion labels from the CheXBERT-annotated impression section of each radiology report, filtering out uncertain and unmentioned labels, and applying a patient-level train/validation/test split to prevent data leakage across sets.

Dataset Source:
https://stanfordaimi.azurewebsites.net/datasets/5158c524-d3ab-4e02-96e9-6ee9efc110a1

Dataset Storage:
https://chexpertdlbenetta.blob.core.windows.net/chexpert-plus?sp=rl&st=2026-05-01T09:36:52Z&se=2026-07-31T17:51:52Z&spr=https&sv=2025-11-05&sr=c&sig=nOL%2Bp4gGEGwsDTXGSuC66jrkEIsoytCgIf37Ep6%2F4Sw%3D

(If the link doesn't work, you should have an invitation in your inbox to access the storage directly)